NFL Showdown DFS: What Are the Sims Saying?

CODY MAIN // SEP 4, 2026

Earlier this offseason, **Mike Leone** published a data-driven piece on The Solver's NFL Classic sims, providing a meticulous deep dive into whether the post-lock simulations actually predict real-world outcomes. Using two seasons' worth of Game Changer data, Leone found that the sims were well-calibrated; lineups that sim positively at lock cash more often, finish in the top 10 more often, and produce better actual ROI than their negatively-simming counterparts. Leone was able to separate the signal from DFS noise and find actionable takeaways from the post-lock sims that we can all apply in our pre-lock lineup construction process.

I entered the doldrums of the NFL offseason with a similar goal, applied to a much different dataset. The question was simple: Do the sims work the same way in NFL showdown as they do in classic?

To find out, I gathered all available post-lock sims data from The Solver's showdown contests, specifically focused on both the Wildcat ($333 buy-in) and Field General ($444 buy-in), for every week we had that data and ETR projection data. The resulting dataset: 49,146 total lineups across 33 different slates, from Week 5 through the Super Bowl. For every lineup reviewed, we can see all of the relevant post-lock sim data (ROI, cash rate, top-1% rate, etc.), along with the actual contest result (lineup rank, money won) and a full construction breakdown (CPT position, roster construction, salary used, duplicates, projected ownership, etc.) thanks to the ETR projection data.

There are a few important distinctions worth pointing out: While this dataset covers 33 different slates, that is still a remarkably small sample. Additionally, this is a very different contest than the one Leone covered, featuring a larger field, lower buy-in, and multi-entry vs. single-entry status. Some of the things Leone unearthed through his review of the classic post-lock sims will hold up for showdown, and others won't. We'll start by asking whether the sims are actually calibrated, what that means for our bankroll, and wrap up by identifying practical construction principles we can apply once the 2026 season gets underway.

Before we dive specifically into our showdown post-lock sims, it's important to review Leone's article to understand what our sims *do* and *don't* account for. With that serving as a baseline for our analysis, let's start by determining whether or not the sims correctly distribute the cash probability across individual lineups. If the sims say a lineup has a 30% cash rate, does that lineup actually cash more often than one the sims have at 15%?

We bucketed all 49,146 lineups by their sim cash rate and compared it to how often lineups in that bucket *actually* cashed:

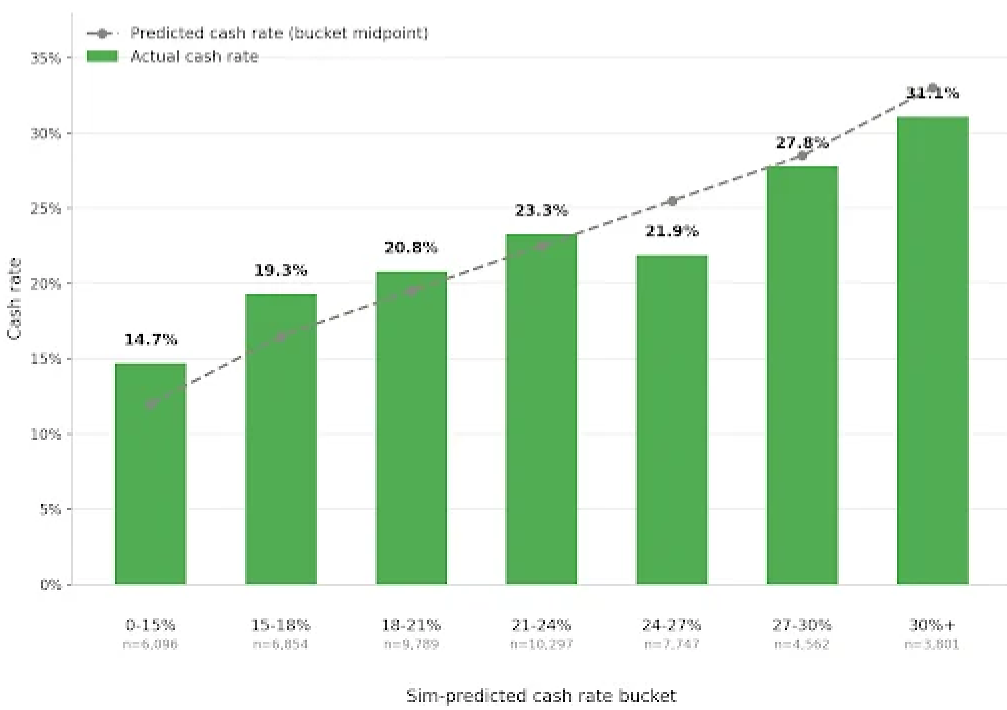

As predicted cash rate increased, actual cash rate increased as well. For example, lineups in the 18-21% sim cash rate bucket cashed 20.8% of the time, while lineups that were predicted to cash at a 30+% clip cashed 31.1% of the time. The sims' cash rate is a reliable signal; a lineup predicted to cash more often than another lineup does, in fact, cash more often.

Do the Sims Pick Winners?

A well-calibrated cash rate is great, but for those playing top-heavy GPP contests, squeaking inside the cash line isn't the primary objective. The most intuitive way to evaluate whether or not the sims have any predictive power is to look at winning lineups and ask: How did they sim?

Leone found that 65.7% of Game Changer winners had a positive sim ROI, well above the 40% of overall lineups that simmed positively in those contests. Our showdown data tells a different story: Of the 33 contest winners in our dataset, just 13 (39.4%) had a positive sim ROI at lock, below the 50.5% of all lineups that simmed positively overall.

It is important to reiterate that there are significant differences between the contest Leone studied and this one. The Game Changer is a single-entry contest, where every user built one lineup. Our Wildcat and Field General database is for multi-entry contests, where 16 of the 33 aforementioned winners (48.5%) were players who entered the maximum number of lineups allowed in that contest. When high-volume players are submitting a diverse portfolio of lineups, it's likelier that a poorly simming lineup emerges as a potential winner.

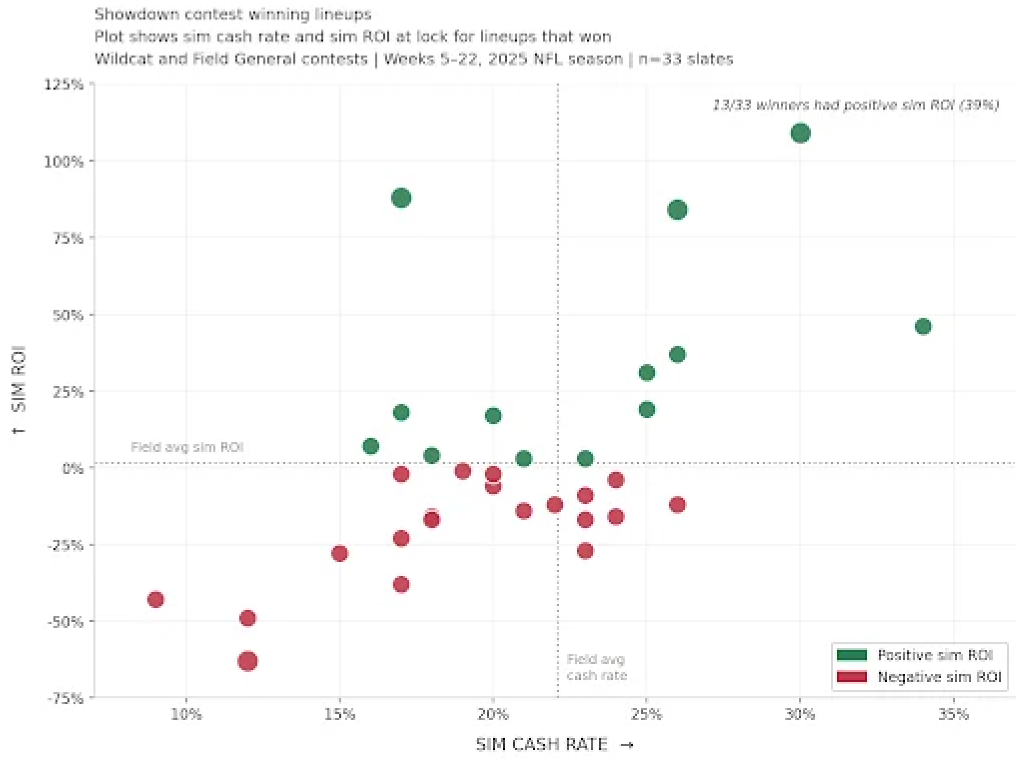

As you can see, winners aren't clustered in the upper-right quadrant quite like they are in the Game Changer, and there are negatively-simming winners across every cash-rate bucket. A few things are working against the sim here: larger fields, multi-entry format, the inherent variance of a single-game slate, *and* the small-sample nature of the data we've collected. Negatively-simming lineups win more often than might be expected, and a portion of those come from high-volume players who were more likely to have a winning lineup in their portfolio, regardless of individual lineup quality. This, by itself, doesn't mean the sim is useless; we've already determined that there is signal in the cash-rate calibration. But the charted winners across this 33-slate sample do provide a useful reality check on what the sims can tell us in this format.

That all felt pretty doomsday-ish as I was typing it: "Sorry, single-game slates are highly volatile, high-volume players are likelier to have a winning lineup in their MME portfolio, and lineups that the sims think are 'bad' at lock win more often than we expect them to."

If we zoom out from first place, the picture changes considerably. Rather than just looking at winning lineups from each contest, let's look at the top-10 finishers across our 33 slates, 375 lineups in total. This gives us both a larger sample of top-finishing lineups and a more meaningful test of what the sims are actually trying to accomplish. It isn't claiming to tell us which single lineup will win a given week, but rather which lineups have the best probability in a top-prize-paying position. Winning a 1,500-person showdown contest still requires a huge amount of positive variance that is tough to reliably predict, especially over such a small sample. But consistently finishing in the top 10, for example, provides a better proxy.

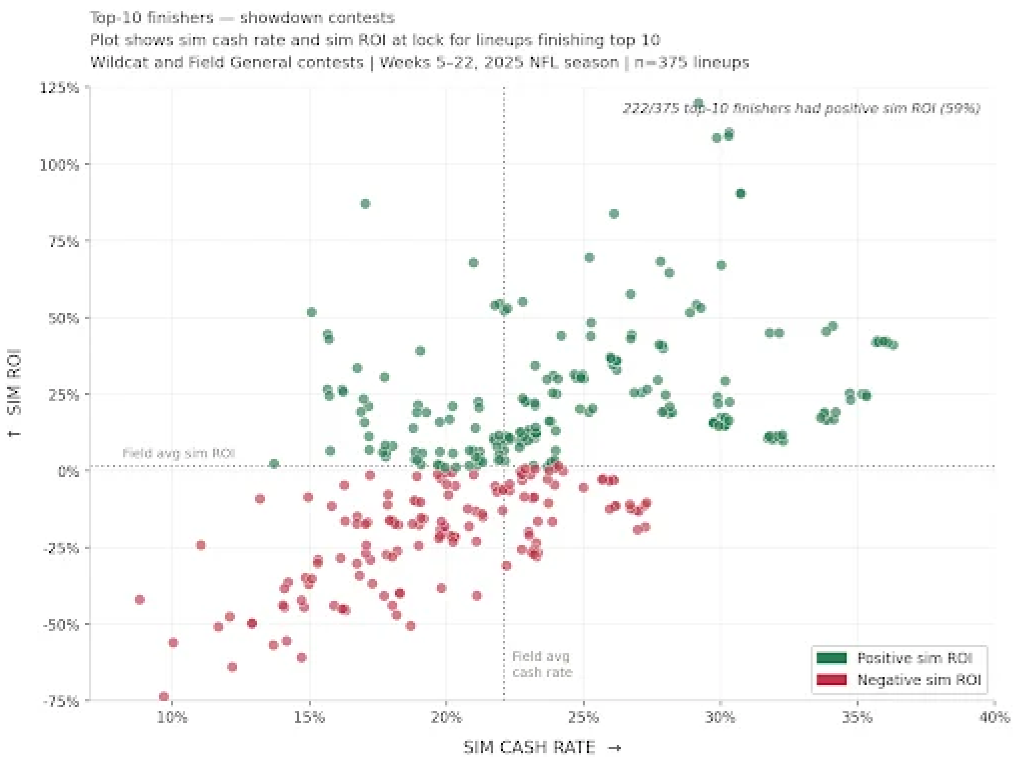

The top-10 chart still has its fair share of negatively-simming lineups, clustered in the lower-left quadrant, but of the 375 top-10 lineups in this dataset, 222 (59.2%) had a positive sim ROI at lock, compared to 50.5% of all lineups in the field. A positively-simming lineup was 1.17x more likely to finish in the top 10 than you'd expect by chance, while negatively-simming lineups were 0.82x as likely. The average sim ROI of top-10 finishers was +6.7%, more than five percentage points above the field average of 1.5%. The sims still struggle to identify exactly which lineup will win on any given showdown slate, but it is providing signal on the pool of lineups more likely to finish where the real money is.

Sim ROI: Signal, Noise, and Dupes

We've already found that the cash-rate calibration shows that the sim reliably separates good lineups from bad ones, and the top-10 rate shows that the sim is correctly identifying lineups that have a chance to finish near the top of GPP leaderboards. However, we still need to explore the relationship between sim ROI and actual ROI and finally touch on the dupe discussion.

We bucketed all 49,031 lineups into eight sim ROI ranges and measured three things: actual ROI, cash rate, and dupe count. Each chart deserves an explanation, and while the findings aren't nearly as clean as we may expect, the takeaways are the most actionable pieces of information from this entire article.

The actual ROI chart isn't the clean staircase you'd expect. The worst lineups according to the post-lock sims are, however, the worst at realizing a positive ROI. Lineups with a -40% or worse post-lock sim ROI posted a -42.0% actual ROI, the worst among all cohorts, while the bucket from -40 to -25% wasn't far behind, checking in with a -24.7% actual ROI. Additionally, the best lineups, those in the 40+% sim ROI group, yielded positive results, averaging an 11.0% actual ROI. With nothing else, this would paint a pretty clear picture: The worst lineups at lock were the least likely to yield profitable results, and the best lineups at lock were. The middle of the distribution, though, raises some eyebrows, with the -25 to -10% bucket outperforming every positive post-lock sim bucket with a 13.8% actual ROI, while every bucket from 0-10%, 10-25%, and 25-40% all produced a negative actual ROI. While Leone found similar variance in these 'middle' buckets in classic, the effect is even more pronounced in showdown.

In the second section of the chart, we can see that cash rate and top-1% finish rate by sim ROI tell a cleaner story. Cash rate rises fairly consistently from the 16.1% in the worst bucket to 24.9% in the 25-40% bucket, while top-1% rate is directionally correct at both extremes as well, checking in at 0.57% in the <= -40% bucket and 1.41% in the >= 40% bucket. So, if the sim is largely doing its job, why the actual ROI dip in the middle buckets that we should expect to outperform their worse-simming counterparts?

The final section highlights average duplicate count for each sim ROI bucket, starting at 1.79 average dupes in the <= -40% range, climbing to a peak of 4.73 in the 0-10% bucket, before leveling off in the >= 40% bucket at 3.89 average dupes. This intuitively makes sense for two reasons: The field is identifying theoretically sound lineups that are more likely to sim well and playing them at a higher rate, and lineups that sim well pre-lock are also more likely to be played by those of us deploying the sims. As we know, when duped lineups finish in the money, winnings are split amongst all lineups, suppressing the actual ROI of that lineup.

We can try to quantify just how much the ROI is suppressed. Among the lineups that actually cashed in the contest, unique entries averaged $1,826 in winnings, for a 396% ROI. Lineups that were deemed “heavily duplicated” (>= 11x) that cashed averaged just $976 in winnings for a 156% ROI, good for a 60% haircut on actual return simply for being duplicated.

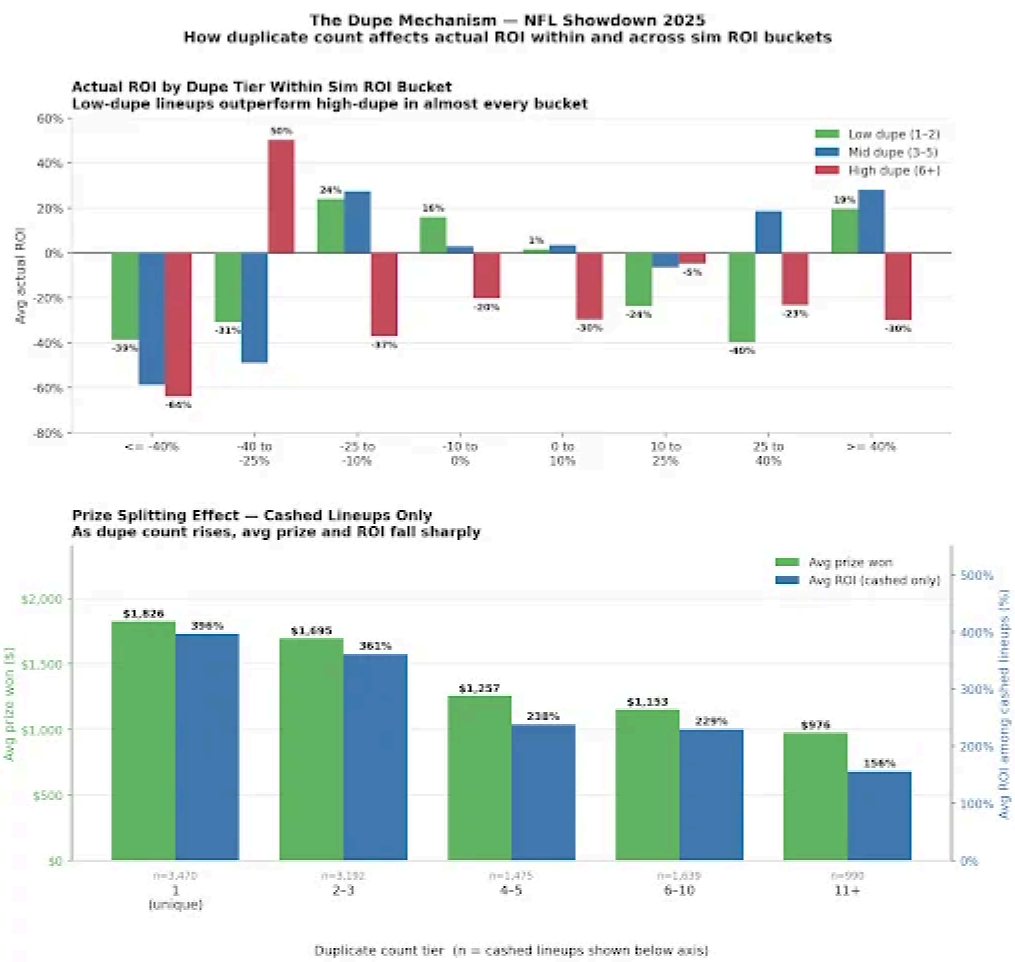

When we split the sim ROI buckets further down into low dupe (1-2), mid dupe (3-5), and high dupe (6+), we can see consistent examples of low-dupe lineups outperforming their higher-duped counterparts in actual ROI. No better example than the >= 40% sim ROI bucket, where low-dupe lineups posted a 19.5% actual ROI compared to -30% for lineups that were duped at least 6x.

It is worth noting that the small 33-slate sample isn't free from some noise here, specifically in the -40 to -25% sim ROI bucket where high-dupe lineups posted a 50% actual ROI. That return is driven by a few lineups, most notably a seven-dupe lineup in Week 18 that actually shipped the slate for an 8,373% ROI. However, 76% of the lineups in the entire dataset from this bucket posted an actual ROI of -100%.

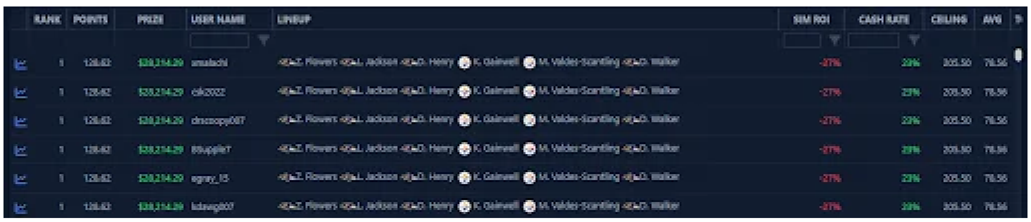

RANK	POINTS	PRIZE	USER NAME	LINEUP	SIM ROI	CASH RATE	CEILING	AVG
1	128.62	$28,214.29	amalachi	Z. Flowers L. Jackson D. Henry K. Gainwell M. Valdes-Scantling D. Walker	-27%	23%	205.50	78.56
1	128.62	$28,214.29	cik2022	Z. Flowers L. Jackson D. Henry K. Gainwell M. Valdes-Scantling D. Walker	-27%	23%	205.50	78.56
1	128.62	$28,214.29	dnccoopy087	Z. Flowers L. Jackson D. Henry K. Gainwell M. Valdes-Scantling D. Walker	-27%	23%	205.50	78.56
1	128.62	$28,214.29	89upple7	Z. Flowers L. Jackson D. Henry K. Gainwell M. Valdes-Scantling D. Walker	-27%	23%	205.50	78.56
1	128.62	$28,214.29	egray_15	Z. Flowers L. Jackson D. Henry K. Gainwell M. Valdes-Scantling D. Walker	-27%	23%	205.50	78.56
1	128.62	$28,214.29	kdawg807	Z. Flowers L. Jackson D. Henry K. Gainwell M. Valdes-Scantling D. Walker	-27%	23%	205.50	78.56

The biggest takeaway from this section is that the worst sim ROI buckets (<= -25%) are genuinely bad lineups and, by and large, should be avoided. These lineups cashed infrequently (16.1%), had the lowest realized return (-42.0%), and did so despite averaging the lowest number of dupes (1.79) of any sim ROI bucket. I'd also add that two lineups with similar sim ROIs may not be equal; the lineup with fewer projected dupes and a 15% sim ROI is likely better than one with 10 dupes and a 20% sim ROI. The biggest problem in deciding between the two is that the actual number of times a lineup is played by our opponents in any given contest is extremely difficult to predict. It's ultimately up to each of us to consider a lineup's desirability and to determine if it's more or less likely to be played than projected.

Contest Selection

When I set out to review post-lock sim data, I knew the Wildcat and Field General-style contests would be a perfect place to start for a few reasons. First, I've only ever dived deep into full-blown lotto-style contests, wager $20 to win $1 million. Second, these mid-tier buy-in contests offer a solid mix of users entering one or just a few lineups and another group max-entering. Regardless, the stakes make it so that *most* people playing in these contests have some level of intentionality behind their lineups. The casual player building a few lineups from the toilet is largely priced out. It's an important distinction because the gap, even in these contests, between single-entry and max-entry players is significant and may provide some insight into how to better deploy our bankroll via proper contest selection.

In the 33 slates we observed, we found that low-entry players (1-3 lineups) accounted for 38.0% of all lineups, mid-entry (4-30) accounted for 35.3%, and max-entry players (31+) accounted for 26.7%. Despite accounting for a smaller portion of the total pool, 31+ entry players accounted for 31.8% of money won compared to just 21.2% for single-entry users.

The post-lock sim data agrees: Max-entry players are simply building better lineups than their single-entry opponents. Single-entry players averaged a -10.5% sim ROI across their lineups, with just 35.7% of their total portfolio simming positively. Players who entered at least 31 lineups averaged a 12.0% sim ROI, with 63.7% of their lineups simming positively. High-volume players weren't just winning on volume; they were winning because they were consistently submitting better lineups.

The actual outcomes align closely with the sim predictions. Single-entry players cashed at a 21.1% clip, found top-1% scores 0.92% of the time, and averaged a -21.6% actual ROI. Max-entry players cashed at a 23.1% rate, found top-1% scores 1.37% of the time, and closed with a +8.0% average actual ROI. Single-entry players had 28.5% of their lineups land in the "Horrid" (< -30%) tier and just 6.9% land in the "Very Good" (>= 40%) tier. Among the 31+ entry players, only 9.3% of lineups landed in the "Horrid" tier, while 18.3% found their way into the "Very Good" tier.

Ultimately, the post-lock sim data reinforces some points many of us already knew: If you're considering playing a single lineup in a multi-entry contest like the Wildcat or Field General, those funds are better deployed in separate single-entry contests at lower buy-ins where users do not have as strong of a structural advantage over you and make up a smaller percentage of the total pool.

Conclusion: Applying What the Sims Are Actually Telling Us

Now that we've reviewed 49,031 showdown lineups across 33 slates, we can attempt to answer the original question: What are the sims saying?

We know the cash-rate calibration is the cleanest signal in the entire 33-slate sample. The sims correctly distribute cash-rate probability across individual lineups from the 14.7% actual cash rate in the lowest bucket to 31.1% in the highest bucket. If the sim says your lineup is more likely to cash than another, it is.

Additionally, the top-1% rate calibration holds up at both extremes but is noisier in the middle of the distribution. Lineups with the lowest sim top-1% rate finished in the top 1% at just 0.52x the expected rate, while lineups in the 1.5%+ top-1% bucket finished in the top 1% at a 1.21-1.29% clip. In terms of "getting shots on goal," prioritizing lineups with higher predicted top-1% rates and avoiding those that don't is a good signal for building a portfolio of lineups with paths to climbing DFS leaderboards.

The harshest reality check from this exercise was that the sim only correctly identified contest winners less than 40% of the time; that is, only 13 of 33 winning lineups simmed positively at lock. If we expand to top-10 finishes, the signal flips in the sims' favor: 59.2% of top-10 finishing lineups had a positive sim ROI.

The dupe equation adds a layer of complexity that we rarely have to worry about with classic slates. The sim is correctly finding quality lineups, reflected in the cash rate and top-1% finish rate, but the actual ROI is also impacted by how many other players are duplicating our lineup. This remains the single biggest user-driven advantage in showdown: understanding lineup desirability and the likelihood other users will also be playing said lineup.

Lastly, the ability to subset our data between max-entry and single-entry users highlights an important decision point non-maxers should consider before registering for Wildcat and Field General-style showdown contests. Max-entry users are systematically building better lineups by almost any metric we can track, including sim ROI, cash rate, top-1% rate, and actual ROI. For single- or low-entry players in MME-style contests, it's worth seriously considering if your bankroll could be better deployed in contests that aren't skewed against you.

Now that we have a firm understanding of what the sims are telling us, we'll take everything we learned here and apply it directly to showdown lineup construction as we set our sights on 2026. In Part 2 of this series, we'll cover CPT selection, roster construction, and how to use all of the above information to turn the sims into a powerful lineup-building tool.

In-Season Package, Showdown

SUBSCRIPTIONS

NFL Products
NBA Products
Golf Products
The Solver

ARTICLES

NFL Content Schedule
Silva's Matchups
Best Ball
Free Podcasts

DFS

Projections
Strategy
DFS Optimizer
Simulations

ETR TEAM

About
YouTube
Discord
Twitter